



DD2380 Artificial Intelligence

Lecture 10 Classical Planning

Jana Tumova

Planning

Deliberating a plan of action to achieve one's goal



Real-world planning

- Limited resources: time, cost, capacity,
- Uncertainty
- Multiple agents
- Different criteria: optimality,
- Robotics: dynamical constraints
- Integration into context, integration with other AI methods

Real-world planning

- Limited resources: time, cost, capacity,
- Uncertainty
- Multiple agents
- Different criteria: optimality,
- Robotics: dynamical constraints
- Integration into context, integration with other AI methods

Today's lecture: only fully observable, deterministic, static, environments.

Planning problem

Planning problem:

- Initial state
- Actions available in a state $ACTIONS(s)$
- Results of applying action $RESULT(s, a)$
- Goal test

Planning vs. problem solving

- Problem solving (search and games):
 - Explicit atomic representations
 - Need good domain-specific heuristics
- Classical planning:
 - Factored representation
 - Domain-independent algorithms

Challenges

- How do we represent a planning problem?
- How do we solve a planning problem?

Challenges

- How do we represent a planning problem?
 - PDDL, STRIPS, ...
- How do we solve a planning problem?

Planning Domain Definition Language (PDDL)

A careful balance between expressivity and simplicity

States

- $At(P, SFO)$
- $At(P, Arlanda) \wedge Plane(P) \wedge Loaded(Cargo, P)$

Actions (think of them as universally quantified)

$Action(Fly(p, from, to),$

PRECOND: $At(p, from) \wedge Plane(p) \wedge Airport(from) \wedge Airport(to)$

EFFECT: $\neg At(p, from) \wedge At(p, to)$)

Example: Air Cargo Transport in PDDL

Init($At(C_1, SFO) \wedge At(C_2, JFK) \wedge At(P_1, SFO) \wedge At(P_2, JFK)$
 $\wedge Cargo(C_1) \wedge Cargo(C_2) \wedge Plane(P_1) \wedge Plane(P_2)$
 $\wedge Airport(JFK) \wedge Airport(SFO)$)

Goal($At(C_1, JFK) \wedge At(C_2, SFO)$)

Action(*Load*(*c*, *p*, *a*),

PRECOND: $At(c, a) \wedge At(p, a) \wedge Cargo(c) \wedge Plane(p) \wedge Airport(a)$

EFFECT: $\neg At(c, a) \wedge In(c, p)$)

Action(*Unload*(*c*, *p*, *a*),

PRECOND: $In(c, p) \wedge At(p, a) \wedge Cargo(c) \wedge Plane(p) \wedge Airport(a)$

EFFECT: $At(c, a) \wedge \neg In(c, p)$)

Action(*Fly*(*p*, *from*, *to*),

PRECOND: $At(p, from) \wedge Plane(p) \wedge Airport(from) \wedge Airport(to)$

EFFECT: $\neg At(p, from) \wedge At(p, to)$)

PDDL States

- State is a conjunction of ground, functionless atomic fluents
 - $At(P, Arlanda) \wedge Plane(P) \wedge Loaded(Cargo, P)$, but
 - $\text{Not } \neg At(P, Bromma)$,
 - $\text{Not } At(x, y)$,
 - $\text{Not } At(P, TheHomeAirport(SAS))$
- Database semantics: Fluents that are not mentioned are false
- Two equivalent viewpoints:
 - Conjunction of fluents: logical inference
 - Set of fluents: set operations

PDDL Action Schemes

Actions and their results are represented through action schemes

- Action name with the list of all used variables
- Precondition: a conjunction of literals saying when an action is *applicable* in a state s , namely if s entails the precondition

$$ACTIONS(s) = \{a \mid s \models PRECOND(a)\}.$$

- Effect: a conjunction of literals representing the literals that need to be removed and added

$$RESULT(s, a) = (s \setminus DEL(a)) \cup ADD(a),$$

where $DEL(a)$ are the fluents that appear as negative literals in the effect and $ADD(a)$ are the fluents that appear as positive ones.

Q: What are the initial state and the goal test?




Action(*Slide*(*t*, *s*₁, *s*₂))


PRECOND : *On*(*t*, *s*₁) ∧ *Tile*(*t*) ∧ *Blank*(*s*₂) ∧ *Adjacent*(*s*₁, *s*₂)

EFFECT : *On*(*t*, *s*₂) ∧ *Blank*(*s*₁) ∧ ¬*On*(*t*, *s*₁) ∧ ¬*Blank*(*s*₂)

Block World in PDDL




Block World in PDDL



Init(*On*(*A*, *Table*) \wedge *On*(*B*, *Table*) \wedge *On*(*C*, *A*)
 \wedge *Block*(*A*) \wedge *Block*(*B*) \wedge *Block*(*C*) \wedge *Clear*(*B*) \wedge *Clear*(*C*))
Goal(*On*(*A*, *B*) \wedge *On*(*B*, *C*))
Action(*Move*(*b*, *x*, *y*),
 PRECOND: *On*(*b*, *x*) \wedge *Clear*(*b*) \wedge *Clear*(*y*) \wedge *Block*(*b*) \wedge *Block*(*y*) \wedge
 (*b* \neq *x*) \wedge (*b* \neq *y*) \wedge (*x* \neq *y*),
 EFFECT: *On*(*b*, *y*) \wedge *Clear*(*x*) \wedge \neg *On*(*b*, *x*) \wedge \neg *Clear*(*y*))
Action(*MoveToTable*(*b*, *x*),

Block World in PDDL



Init($On(A, Table) \wedge On(B, Table) \wedge On(C, A)$
 $\wedge Block(A) \wedge Block(B) \wedge Block(C) \wedge Clear(B) \wedge Clear(C)$)
Goal($On(A, B) \wedge On(B, C)$)
Action(*Move*(b, x, y),
PRECOND: $On(b, x) \wedge Clear(b) \wedge Clear(y) \wedge Block(b) \wedge Block(y) \wedge$
 $(b \neq x) \wedge (b \neq y) \wedge (x \neq y)$,
EFFECT: $On(b, y) \wedge Clear(x) \wedge \neg On(b, x) \wedge \neg Clear(y)$)
Action(*MoveToTable*(b, x),
PRECOND: $On(b, x) \wedge Clear(b) \wedge Block(b) \wedge (b \neq x)$,
EFFECT: $On(b, Table) \wedge Clear(x) \wedge \neg On(b, x)$)

Challenges


- How do we represent a planning problem?
 - PDDL, STRIPS, ...
- How do we solve a planning problem?

Challenges

- How do we represent a planning problem?
 - PDDL, STRIPS, ...
- How do we solve a planning problem?
 - Via forward search and backward search with heuristics


PDDL Planning Problem as a State Space Search

- Description of a planning problem defines a search problem
- States are truth assignments to fluents, actions and results define the transitions



PDDL Planning Problem as a State Space Search

- Description of a planning problem defines a search problem
- States are truth assignments to fluents, actions and results define the transitions



Forward (progression) search

- As expected
 - Start at the initial state
 - Explore applicable actions
- Properties
 - Large branching factor, often explores irrelevant actions
 - Needs a good heuristic, ideally domain-independent

Example: Air Cargo Transport in PDDL

Init($At(C_1, SFO) \wedge At(C_2, JFK) \wedge At(P_1, SFO) \wedge At(P_2, JFK)$
 $\wedge Cargo(C_1) \wedge Cargo(C_2) \wedge Plane(P_1) \wedge Plane(P_2)$
 $\wedge Airport(JFK) \wedge Airport(SFO)$)

Goal($At(C_1, JFK) \wedge At(C_2, SFO)$)

Action(*Load*(c, p, a),

PRECOND: $At(c, a) \wedge At(p, a) \wedge Cargo(c) \wedge Plane(p) \wedge Airport(a)$

EFFECT: $\neg At(c, a) \wedge In(c, p)$)

Action(*Unload*(c, p, a),

PRECOND: $In(c, p) \wedge At(p, a) \wedge Cargo(c) \wedge Plane(p) \wedge Airport(a)$

EFFECT: $At(c, a) \wedge \neg In(c, p)$)

Action(*Fly*($p, from, to$),

PRECOND: $At(p, from) \wedge Plane(p) \wedge Airport(from) \wedge Airport(to)$

EFFECT: $\neg At(p, from) \wedge At(p, to)$)

Backward (regression) search

Suitable, since the predecessor function can be easily derived:

$$g' = (g \setminus ADD(a)) \cup PRECOND(a)$$

Idea: Reduce branching by considering only *relevant* (ground) actions, those that make some literal in the goal true, but don't make any literal in the goal false.

Q: What are the relevant ground actions?

$Init(At(C_1, SFO) \wedge At(C_2, JFK) \wedge At(P_1, SFO) \wedge At(P_2, JFK)$
 $\wedge Cargo(C_1) \wedge Cargo(C_2) \wedge Plane(P_1) \wedge Plane(P_2)$
 $\wedge Airport(JFK) \wedge Airport(SFO))$
 $Goal(At(C_1, JFK) \wedge At(C_2, SFO))$
 $Action(Load(c, p, a),$
 PRECOND: $At(c, a) \wedge At(p, a) \wedge Cargo(c) \wedge Plane(p) \wedge Airport(a)$
 EFFECT: $\neg At(c, a) \wedge In(c, p))$
 $Action(Unload(c, p, a),$
 PRECOND: $In(c, p) \wedge At(p, a) \wedge Cargo(c) \wedge Plane(p) \wedge Airport(a)$
 EFFECT: $At(c, a) \wedge \neg In(c, p))$
 $Action(Fly(p, from, to),$
 PRECOND: $At(p, from) \wedge Plane(p) \wedge Airport(from) \wedge Airport(to)$
 EFFECT: $\neg At(p, from) \wedge At(p, to))$

Q: What are the relevant ground actions?

$Init(At(C_1, SFO) \wedge At(C_2, JFK) \wedge At(P_1, SFO) \wedge At(P_2, JFK)$
 $\wedge Cargo(C_1) \wedge Cargo(C_2) \wedge Plane(P_1) \wedge Plane(P_2)$
 $\wedge Airport(JFK) \wedge Airport(SFO))$
 $Goal(At(C_1, JFK) \wedge At(C_2, SFO))$
 $Action(Load(c, p, a),$
 PRECOND: $At(c, a) \wedge At(p, a) \wedge Cargo(c) \wedge Plane(p) \wedge Airport(a)$
 EFFECT: $\neg At(c, a) \wedge In(c, p))$
 $Action(Unload(c, p, a),$
 PRECOND: $In(c, p) \wedge At(p, a) \wedge Cargo(c) \wedge Plane(p) \wedge Airport(a)$
 EFFECT: $At(c, a) \wedge \neg In(c, p))$
 $Action(Fly(p, from, to),$
 PRECOND: $At(p, from) \wedge Plane(p) \wedge Airport(from) \wedge Airport(to)$
 EFFECT: $\neg At(p, from) \wedge At(p, to))$

Relevant action $Unload(C_2, p', SFO)$

Q: What is the new goal?

$Init(At(C_1, SFO) \wedge At(C_2, JFK) \wedge At(P_1, SFO) \wedge At(P_2, JFK)$
 $\wedge Cargo(C_1) \wedge Cargo(C_2) \wedge Plane(P_1) \wedge Plane(P_2)$
 $\wedge Airport(JFK) \wedge Airport(SFO))$

$Goal(At(C_1, JFK) \wedge At(C_2, SFO))$

$Action(Load(c, p, a),$

PRECOND: $At(c, a) \wedge At(p, a) \wedge Cargo(c) \wedge Plane(p) \wedge Airport(a)$

EFFECT: $\neg At(c, a) \wedge In(c, p)$)

$Action(Unload(c, p, a),$

PRECOND: $In(c, p) \wedge At(p, a) \wedge Cargo(c) \wedge Plane(p) \wedge Airport(a)$

EFFECT: $At(c, a) \wedge \neg In(c, p)$)

$Action(Fly(p, from, to),$

PRECOND: $At(p, from) \wedge Plane(p) \wedge Airport(from) \wedge Airport(to)$

EFFECT: $\neg At(p, from) \wedge At(p, to)$)

Q: What is the new goal?

$Init(At(C_1, SFO) \wedge At(C_2, JFK) \wedge At(P_1, SFO) \wedge At(P_2, JFK)$
 $\wedge Cargo(C_1) \wedge Cargo(C_2) \wedge Plane(P_1) \wedge Plane(P_2)$
 $\wedge Airport(JFK) \wedge Airport(SFO))$
 $Goal(At(C_1, JFK) \wedge At(C_2, SFO))$
 $Action(Load(c, p, a),$
 PRECOND: $At(c, a) \wedge At(p, a) \wedge Cargo(c) \wedge Plane(p) \wedge Airport(a)$
 EFFECT: $\neg At(c, a) \wedge In(c, p))$
 $Action(Unload(c, p, a),$
 PRECOND: $In(c, p) \wedge At(p, a) \wedge Cargo(c) \wedge Plane(p) \wedge Airport(a)$
 EFFECT: $At(c, a) \wedge \neg In(c, p))$
 $Action(Fly(p, from, to),$
 PRECOND: $At(p, from) \wedge Plane(p) \wedge Airport(from) \wedge Airport(to)$
 EFFECT: $\neg At(p, from) \wedge At(p, to))$

Relevant action $Unload(C_2, p', SFO)$

$g' = At(C_1, JFK) \wedge$

$In(C_2, p') \wedge At(p', C_2) \wedge Cargo(C_2) \wedge Plane(p') \wedge Airport(SFO)$

Heuristics

In both progression and regression search, we need good a heuristic.

Domain-independent heuristics

- Any planning problem instance
- Define a relaxed, easier problem and a heuristic as a solution to this easier problem
- Use the internal structure of a factored representation of the state space
- Ignore preconditions
- Ignore delete lists
- State abstraction
- Decomposition
- Planning graph

Domain-independent heuristics

- Any planning problem instance
- Define a relaxed, easier problem and a heuristic as a solution to this easier problem
- Use the internal structure of a factored representation of the state space
- Ignore preconditions
- Ignore delete lists
- State abstraction
- Decomposition
- Planning graph

Ignore preconditions

Ignore all preconditions

- Every action becomes applicable in every state
- The number of edges in the graph increases
- The number of steps to solve the relaxed problem is almost the number of unsatisfied fluents in the goal
 - Some actions may achieve multiple goals
 - Some actions may undo the effects of others

Ignore some preconditions

Example: Ignore preconditions

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

$Action(Slide(t, s_1, s_2))$

$PRECOND : On(t, s_1) \wedge Tile(t) \wedge Blank(s_2) \wedge Adjacent(s_1, s_2)$

$EFFECT : On(t, s_2) \wedge Blank(s_1) \wedge \neg On(t, s_1) \wedge \neg Blank(s_2)$

$h_1(n)$: number of the misplaced tiles

Example: Ignore preconditions



$Action(Slide(t, s_1, s_2))$


$PRECOND : On(t, s_1) \wedge Tile(t) \wedge Blank(s_2) \wedge Adjacent(s_1, s_2)$

$EFFECT : On(t, s_2) \wedge Blank(s_1) \wedge \neg On(t, s_1) \wedge \neg Blank(s_2)$

$h_1(n)$: number of the misplaced tiles

- The relaxed problem assumed we can transfer any tile anywhere
- $PRECOND : On(t, s_1) \wedge Tile(t) \wedge \neg Blank(s_2) \wedge \neg Adjacent(s_1, s_2)$

Q: Match ignoring preconditions with the heuristic




Action(*Slide*(*t*, *s*₁, *s*₂))

PRECOND : *On*(*t*, *s*₁) ∧ *Tile*(*t*) ∧ *Blank*(*s*₂) ∧ *Adjacent*(*s*₁, *s*₂)

EFFECT : *On*(*t*, *s*₂) ∧ *Blank*(*s*₁) ∧ ¬*On*(*t*, *s*₁) ∧ ¬*Blank*(*s*₂)

*h*₂(*n*): sum of the distances of the tiles from the goal position

Q: Match ignoring preconditions with the heuristic



Action(*Slide*(t, s_1, s_2))

PRECOND : $On(t, s_1) \wedge Tile(t) \wedge Blank(s_2) \wedge Adjacent(s_1, s_2)$

EFFECT : $On(t, s_2) \wedge Blank(s_1) \wedge \neg On(t, s_1) \wedge \neg Blank(s_2)$

$h_2(n)$: sum of the distances of the tiles from the goal position

- The relaxed problem assumed we can transfer any tile to an adjacent cell

PRECOND : $On(t, s_1) \wedge Tile(t) \wedge \neg Blank(s_2) \wedge Adjacent(s_1, s_2)$

Domain-independent heuristics

- Any planning problem instance
- Define a relaxed, easier problem and a heuristic as a solution to this easier problem
- Use the internal structure of a factored representation of the state space
- Ignore preconditions
- Ignore delete lists
- State abstraction
- Decomposition
- Planning graph

Ignore Delete Lists

- Remove all negative literals from effects
- No action will undo progress by another action towards the goal
- The number of edges in the graph increases

Domain-independent Heuristics

- Any planning problem instance
- Define a relaxed, easier problem and a heuristic as a solution to this easier problem
- Use the internal structure of a factored representation of the state space
- Ignore preconditions
- Ignore delete lists
- State abstraction
- Decomposition
- Planning graph

State Abstraction

- Many-to-one mapping from states in the ground representation to the abstract representation
- The number of states decreases
- Ignore some fluents

Domain-independent Heuristics

- Any planning problem instance
- Define a relaxed, easier problem and a heuristic as a solution to this easier problem
- Use the internal structure of a factored representation of the state space
- Ignore preconditions
- Ignore delete lists
- State abstraction
- Decomposition
- Planning graph

Decomposition

- $G = G_1 \wedge \dots \wedge G_n$
- Instead of problem P , we solve problems P_1, \dots, P_n
- We can use $\max_i \text{COST}(P_i)$ as a heuristic

Decomposition

- $G = G_1 \wedge \dots \wedge G_n$
- Instead of problem P , we solve problems P_1, \dots, P_n
- We can use $\max_i \text{COST}(P_i)$ as a heuristic
- Q: Can we use $\text{COST}(P_1) + \dots + \text{COST}(P_n)$ as a heuristic?

Decomposition

- $G = G_1 \wedge \dots \wedge G_n$
- Instead of problem P , we solve problems P_1, \dots, P_n
- We can use $\max_i \text{COST}(P_i)$ as a heuristic
- Q: Can we use $\text{COST}(P_1) + \dots + \text{COST}(P_n)$ as a heuristic?
- Q: What generally happens if we use a non-admissible heuristic?

Domain-independent Heuristics

- Any planning problem instance
- Define a relaxed, easier problem and a heuristic as a solution to this easier problem
- Use the internal structure of a factored representation of the state space
- Ignore preconditions
- Ignore delete lists
- State abstraction
- Decomposition
- Planning graph

Planning Graph

Init(*Have*(*Cake*))

Goal(*Have*(*Cake*) \wedge *Eaten*(*Cake*))

Action(*Eat*(*Cake*))

PRECOND: *Have*(*Cake*)

EFFECT: \neg *Have*(*Cake*) \wedge *Eaten*(*Cake*)

Action(*Bake*(*Cake*))

PRECOND: \neg *Have*(*Cake*)

EFFECT: *Have*(*Cake*)

Planning Graph

Init(*Have*(*Cake*))
Goal(*Have*(*Cake*) \wedge *Eaten*(*Cake*))
Action(*Eat*(*Cake*))
 PRECOND: *Have*(*Cake*)
 EFFECT: \neg *Have*(*Cake*) \wedge *Eaten*(*Cake*)
Action(*Bake*(*Cake*))
 PRECOND: \neg *Have*(*Cake*)
 EFFECT: *Have*(*Cake*)

State level

S_0

Have(*Cake*)

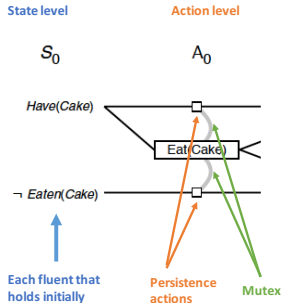
\neg *Eaten*(*Cake*)



Each fluent that
holds initially

Planning Graph

$Init(Have(Cake))$
 $Goal(Have(Cake) \wedge Eaten(Cake))$
 $Action(Eat(Cake))$
 PRECOND: $Have(Cake)$
 EFFECT: $\neg Have(Cake) \wedge Eaten(Cake)$
 $Action(Bake(Cake))$
 PRECOND: $\neg Have(Cake)$
 EFFECT: $Have(Cake)$

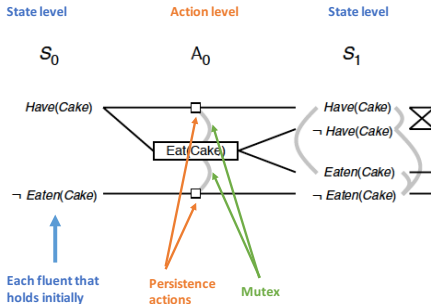


Mutex:

- Inconsistent effects: one action negates the effect of the other
- Interference: one of the effects of one action is the negation of a precondition of the other
- Competing needs: one of the preconditions of one action is the negation of a precondition of the other

Planning Graph

$Init(Have(Cake))$
 $Goal(Have(Cake) \wedge Eaten(Cake))$
 $Action(Eat(Cake))$
 PRECOND: $Have(Cake)$
 EFFECT: $\neg Have(Cake) \wedge Eaten(Cake)$
 $Action(Bake(Cake))$
 PRECOND: $\neg Have(Cake)$
 EFFECT: $Have(Cake)$

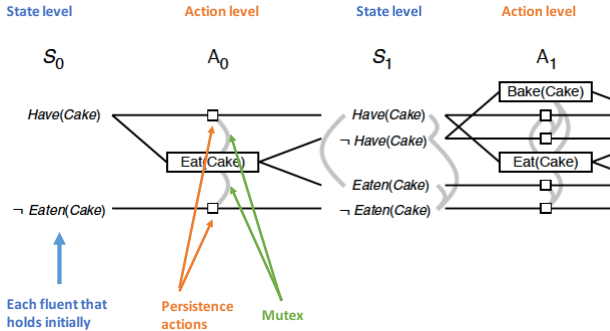


Mutex for literals:

- Negation
- If each possible pair of actions that could achieve the two literals is mutually exclusive


Planning Graph

$Init(Have(Cake))$
 $Goal(Have(Cake) \wedge Eaten(Cake))$
 $Action(Eat(Cake))$
 PRECOND: $Have(Cake)$
 EFFECT: $\neg Have(Cake) \wedge Eaten(Cake)$
 $Action(Bake(Cake))$
 PRECOND: $\neg Have(Cake)$
 EFFECT: $Have(Cake)$



Planning Graph

$Init(Have(Cake))$
 $Goal(Have(Cake) \wedge Eaten(Cake))$
 $Action(Eat(Cake))$
 PRECOND: $Have(Cake)$
 EFFECT: $\neg Have(Cake) \wedge Eaten(Cake)$
 $Action(Bake(Cake))$
 PRECOND: $\neg Have(Cake)$
 EFFECT: $Have(Cake)$



Planning Graph

- Directed graph with alternating state and action levels
- Roughly,
 - S_i level contains literals that could hold at time i
 - A_i level contains actions that could have their preconditions satisfied at time i .
- Persistence actions: a literal can persist if no action negates it
- Mutex:
 - Inconsistent effects: one action negates effect of the other
 - Interference: effect of one action negates a precondition of the other
 - Competing needs: the precondition of one action is mutually exclusive with a precondition of the other

Planning Graph Properties

- Polynomial in the size of the planning problem
- A literal never appears too late, but might appear too early
- If a goal literal does not appear in the final level of the graph, the problem is unsolvable
- We can use it for designing an independent-domain heuristic:
 - Max-level heuristic: admissible, but not always accurate
 - Level-sum heuristic: generally inadmissible, but works well in practice
 - Set-level heuristic: find the level at which all the goal literals appear without being mutually exclusive; admissible and works well

Challenges

- How do we represent a planning problem?
 - PDDL, STRIPS, ...
- How do we solve a planning problem?
 - Via forward search and backward search with heuristics

Challenges

- How do we represent a planning problem?
 - PDDL, STRIPS, ...
- How do we solve a planning problem?
 - Via forward search and backward search with heuristics
 - Via GRAPHPlan

GRAPHPlan

- Using the planning graph to extract a plan directly instead of using it to design a heuristic for search
- EXTRACT-SOLUTION either through CSP or through backward search in the planning graph, not in the state space
- If EXTRACT-SOLUTION does not find a plan, we store $(level, goals)$ in *nogoods*

function GRAPHPLAN(*problem*) **returns** solution or failure

graph \leftarrow INITIAL-PLANNING-GRAPH(*problem*)

goals \leftarrow CONJUNCTS(*problem*.GOAL)

nogoods \leftarrow an empty hash table

for *tl* = 0 to ∞ **do**

if *goals* all non-mutex in S_t of *graph* **then**

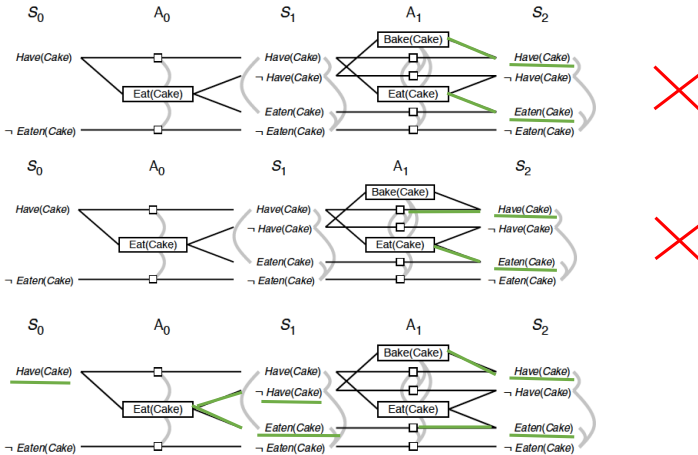
solution \leftarrow EXTRACT-SOLUTION(*graph*, *goals*, NUMLEVELS(*graph*), *nogoods*)

if *solution* \neq failure **then return** *solution*

if *graph* and *nogoods* have both leveled off **then return** failure

graph \leftarrow EXPAND-GRAPH(*graph*, *problem*)

GRAPHPlan



GRAPHPlan Properties

- Literals increase monotonically
 - Actions increase monotonically
 - Mutexes decrease monotonically
 - No-goods decrease monotonically
-
- It is not enough to level-off the graph
 - Termination when mutexes and no-goods have both leveled off

Challenges

- How do we represent a planning problem?
 - PDDL, STRIPS, ...
- How do we solve a planning problem?
 - Via forward search and backward search with heuristics
 - Via GRAPHPlan
 - As refinement of partially ordered plans
 - As Boolean satisfiability
 - As first-order logical deduction: situation calculus
 - As constraint satisfaction

Think-pair-share

- What is the point of using planning as opposed to problem solving?
- What are the limitations of PDDL as we saw it today?

Final remarks

Where to learn more

- Artificial Intelligence: A Modern Approach by Stuart J. Russell and Peter Norvig, chapter 10 + the udacity course Intro to AI
- Automated Planning and Acting by Dana S. Nau, Malik Ghallab, and Paolo Traverso

Tools

- <http://www.fast-downward.org>